

CS528

Multi-threading and OpenMP

A Sahu
Dept of CSE, IIT Guwahati

Outline

- Thread Safety : 4 Classic cases
- Thread Synchronization: Basic
- **Implicit/Auto Thread Pooling: OpenMP/Cilk**

Shared Variables in Threaded C Programs

- Question: Which variables in a threaded C program are shared variables?
 - Answer not as simple as “global variables are shared” and “stack variables are private”
- Requires answers to the following questions:
 - What is the memory model for threads?
 - How are variables mapped to memory instances?

Parallel Counter: without Lock

```
#define NITERS 100
int cnt = 0; /* shared */
int main() {
    pthread_t tid1, tid2;
    pthread_create(&tid1, NULL, count, NULL);
    pthread_create(&tid2, NULL, count, NULL);
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);
    if (cnt != NITERS*2) printf("BOOM! cnt=%d", cnt);
    else printf("OK cnt=%d\n", cnt);
}
```

```
void *count(void *arg) {
for (int i=0; i<NITERS; i++) cnt++;
}
```

```
$/badcnt
BOOM! cnt=196
$/badcnt
BOOM! cnt=184
```

cnt should be 200
What went wrong?!

Thread Safety

- Functions called from a thread must be *thread-safe*
- There are four (non-disjoint) classes of thread-unsafe functions:
 - **Class 1: Failing to protect shared variables : L/UL**
 - Class 2: Relying on persistent state across invocations
 - Class 3: Returning pointer to static variable
 - Class 4: Calling thread-unsafe functions

Class 1: Failing to protect shared variables

- Fix: Use Lock and unlock semaphore operations
- Issue: Synchronization operations will slow down code
- Example: `goodcnt.c`

```
void *count(void *arg) {  
    for(int i=0; i<NITERS; i++)  
        pthread_mutex_lock(&LV);  
        cnt++;  
        pthread_mutex_unlock(&LV);  
} // LV is lock variable
```

Class 2: Relying on persistent state across multiple function invocations

- Random number generator relies on static state
- Fix: Rewrite function so that caller passes in all necessary state, → Maintain Thread Specific State

```
int rand() {  
    static uint next = 1;  
    next = next*1103515245 + 12345;  
    return (uint) (next/65536) % 32768;  
}  
  
void srand(uint seed) {  
    next = seed;  
}
```

Class 3: Returning pointer to **static variable**

- Fixes: 1. Rewrite code so caller passes pointer to `struct`, Issue: Requires changes in caller and callee
- *Lock-and-copy*: Issue: Requires only simple changes in caller (and none in callee), However, caller must free memory

```
struct hostent *gethostbyname (  
                                char* name) {  
    static struct hostent h;  
    <contact DNS and fill in h>  
    return &h;  
}
```


Class 3: Returning pointer to static variable

```
struct hostent *gethostbyname_ts(char *p) {  
    struct hostent *q = Malloc(...);  
    P(&mutex); /* lock */  
    p = gethostbyname(name);  
    *q = *p;    /* copy */  
    V(&mutex);  
    return q;  
}
```

```
hostp = malloc(...);  
gethostbyname_r(name, hostp);
```

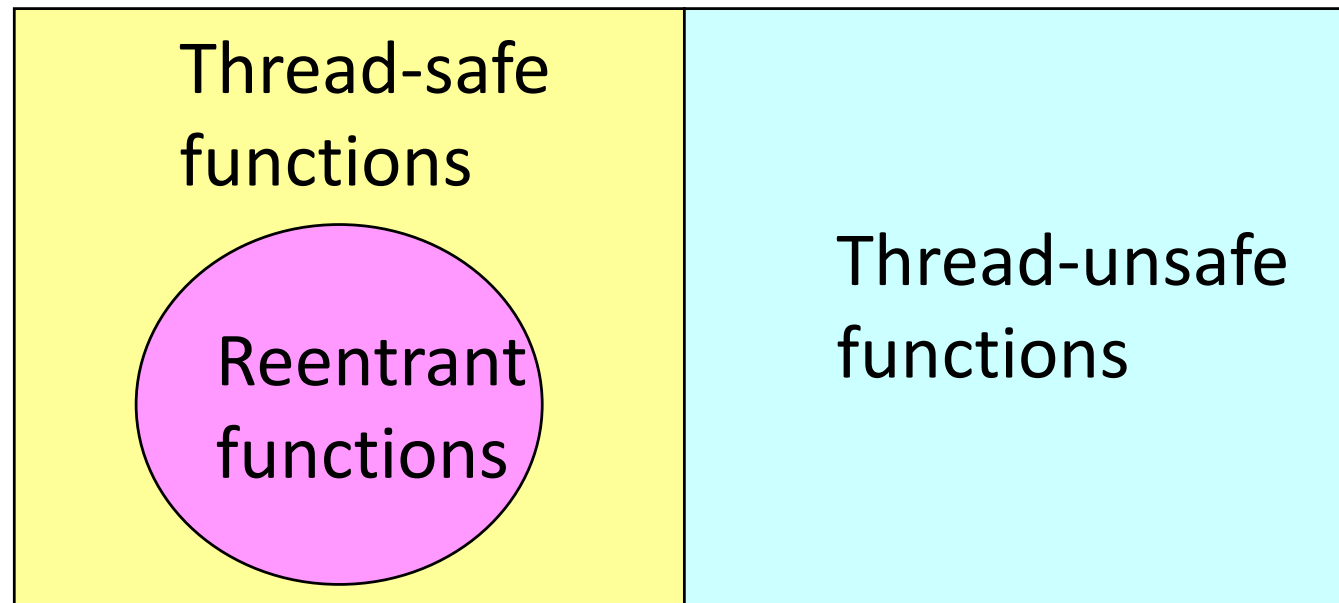
Class 4: Calling thread-unsafe functions

- Calling one thread-unsafe function makes an entire function thread-unsafe
- Fix: Modify the function so it calls only thread-safe functions

Reentrant Functions

- A function is *reentrant* iff it accesses NO shared variables when called from multiple threads
 - Reentrant functions are a proper subset of the set of thread-safe functions
 - NOTE: The fixes to Class 2 and 3 thread-unsafe functions require modifying the function to make it reentrant (only first fix for Class 3 is reentrant)

All functions




Thread-Safe Library Functions

- Most functions in the Standard C Library (at the back of your K&R text) are thread-safe
 - Examples: **malloc**, **free**, **printf**, **scanf**
- All Unix system calls are thread-safe
- Library calls that aren't thread-safe:

Thread-unsafe function	Class	Reentrant version
asctime	3	asctime_r
ctime	3	ctime_r
gethostbyaddr	3	gethostbyaddr_r
gethostbyname	3	gethostbyname_r
inet_ntoa	3	(none)
localtime	3	localtime_r
rand	2	rand_r

Many threads trying to acquire Mutex LOCK



Synchronization Primitives


- **int** pthread_mutex_init(
pthread_mutex_t *mutex_lock,
const pthread_mutexattr_t *lock_attr);
- **int** pthread_mutex_lock(
pthread_mutex_t *mutex_lock);
- **int** pthread_mutex_unlock(
pthread_mutex_t *mutex_lock);
- **int** pthread_mutex_trylock(
pthread_mutex_t *mutex_lock);

Locking Overhead

- Serialization points
 - Minimize the size of critical sections
 - Be careful
- Rather than wait, check if lock is available
 - **pthread_mutex_trylock**
 - If already locked, will return EBUSY
 - Will require restructuring of code
 - **Suspend self by pthread_yield() Give chance to others**
 - **Suspend self by doing a timed wait..**
 - {1, 1, 1,...}, {1, 2, 3, 4,...}, {1,2,4,8,...},...


Performance of Locking

- **Spinning/busy wait waste time**
- Recall MAC Protocol
 - Non Persistence CSMA protocol
 - Wait random time if medium if busy, then send
- Spin lock with exponential back-off reduces contention
 - Wait k amount of time for 1st attempt
 - Wait $k * c^i$ amount of time for i^{th} attempt



Performance of Locking: Monitor

- Spinning/busy wait waste time
- There should be methods for
 - Queuing the shared accessor
 - Calling the next guy when one guy finishes his work
- Best example is : Barber shop
 - Book using Phone, get call before some time of your turn



Locking in Data Structure

- Concurrent Data Structure
 - CDS-List, CSE-Stack, CDS-Queue
 - CDS-PQ, CDS-Hash, CSD-Heap
- Given a Bank with 10^6 account
 - It is not worth to lock whole DB of bank
 - The bank DB may be maintained by List/Hash
 - Suppose A transfer money to B, It is preferred to lock A and Lock B to do the transaction (A.bal-X) and (B.bal+X)
 - Fine grain locking
- Book: Wait free and lock free DS
 - C++ in action : Mileswky
 - Art of multiprocessor programming, by Shavit & Herlihy:
Godel prize winner

Synch. Primitives

- `pthread_mutex_init, lock, unlock, trylock`
- `pthread_cond_wait, signal, broadcast, init, destroy`

Condition Variables for Synchronization

- Condition variable allows a thread
 - To block itself until specified data reaches a predefined state.
- Condition variable
 - Associated with a predicate (P)
 - When the predicate becomes true, the condition variable is used to signal one or more threads waiting on the condition.
- Single condition variable
 - May be associated with more than one predicate
 - Ex: $P = X \text{ OR } Y \text{ AND } (Z \text{ OR } K)$

Synchronization Hierarchy ☺ ☺ ☺

- One == (used by) == > other
- **LL+SC ==> TAS/CAS/FAI/XCHG==>Lock/Unlock**
 - All **TAS/CAS/GAS/FAI/XCHG** do the same work
- Lock/Unlock == > Mutex //Mutex use L/UL
- Mutex == > Semaphore // Semaphore uses Mutex
 - Wait() and Signal()
- Semaphore == > Monitor //Monitor uses Semaphore
 - Many wait/Many Signal, Processes in Queue
 - Monitor : Another Abstract Type
 - *which use semaphore, mutex, conditions*

Lock Variable in Multithreaded App

- Lock Variable (LV): In single core machine looks fine
- **Multithreaded App: Run on Multicore**
- Private Caches: available on Multiple Cores
- **Can the Lock variable be cached?**
 - Lock variable will be at private cache different cores
 - Who will maintain the consistency?
 - Working of Lock variable require: Atomic Transaction
 - Ans: Cache coherence protocol responsible to maintain state of lock variable

C++ Thread:atomic

```
atomic_uint AtomicCount;
```

```
void DoCount () {
```

```
    int j, timesperthrd;
```

```
    timesperthrd=(TIMES/NUM_THREADS);
```

```
    for (j=0; j<timesperthrd; j++) AtomicCount++;
```

```
}
```

```
main () {
```

```
    thread T[N_THRDS]; int i;
```

```
    for (i=0; i<N_THRDS; i++) T[i]=thread(DoCount);
```

```
    for (i=0; i<N_THRDS; i++) T[i].join();
```

```
}
```

Improved

```
atomic_uint AtomicCount;
```

```
void DoCount () {
```

```
    int j, timesperthrd, localcount=0;
```

```
    timesperthrd=(TIMES/NUM_THREADS);
```

```
    for (j=0; j<timesperthrd; j++) localcount++;
```

```
        AtomicCount+=localcount;
```

```
}
```

```
main () {
```

```
    thread T[N_THRDS]; int i;
```

```
    for (i=0; i<N_THRDS; i++) T[i]=thread(DoCount);
```

```
    for (i=0; i<N_THRDS; i++) T[i].join();
```

```
}
```


OpenMP

OpenMP

- Compiler directive: Automatic parallelization
- Auto generate thread and get synchronized

```
#include <openmp.h>
main() {
    #pragma omp parallel
    #pragma omp for schedule(static)
    {
        for (int i=0; i<N; i++) {
            a[i]=b[i]+c[i];
        }
    }
}
```

```
$ gcc -fopenmp test.c
```

```
$ export OMP_NUM_THREADS=4
```

```
$/a.out
```

OpenMP: Parallelism

Sequential code

```
for (int i=0; i<N; i++)  
    a[i]=b[i]+c[i];
```

OpenMP: Parallelism

(Semi) manual parallel

```
#pragma omp parallel
{
    int id =omp_get_thread_num();
    int Nthr=omp_get_num_threads();
    int istart = id*N/Nthr
    int iend= (id+1)*N/Nthr;
    for (int i=istart;i<iend;i++) {
        a[i]=b[i]+c[i];
    }
}
```

OpenMP: Parallelism

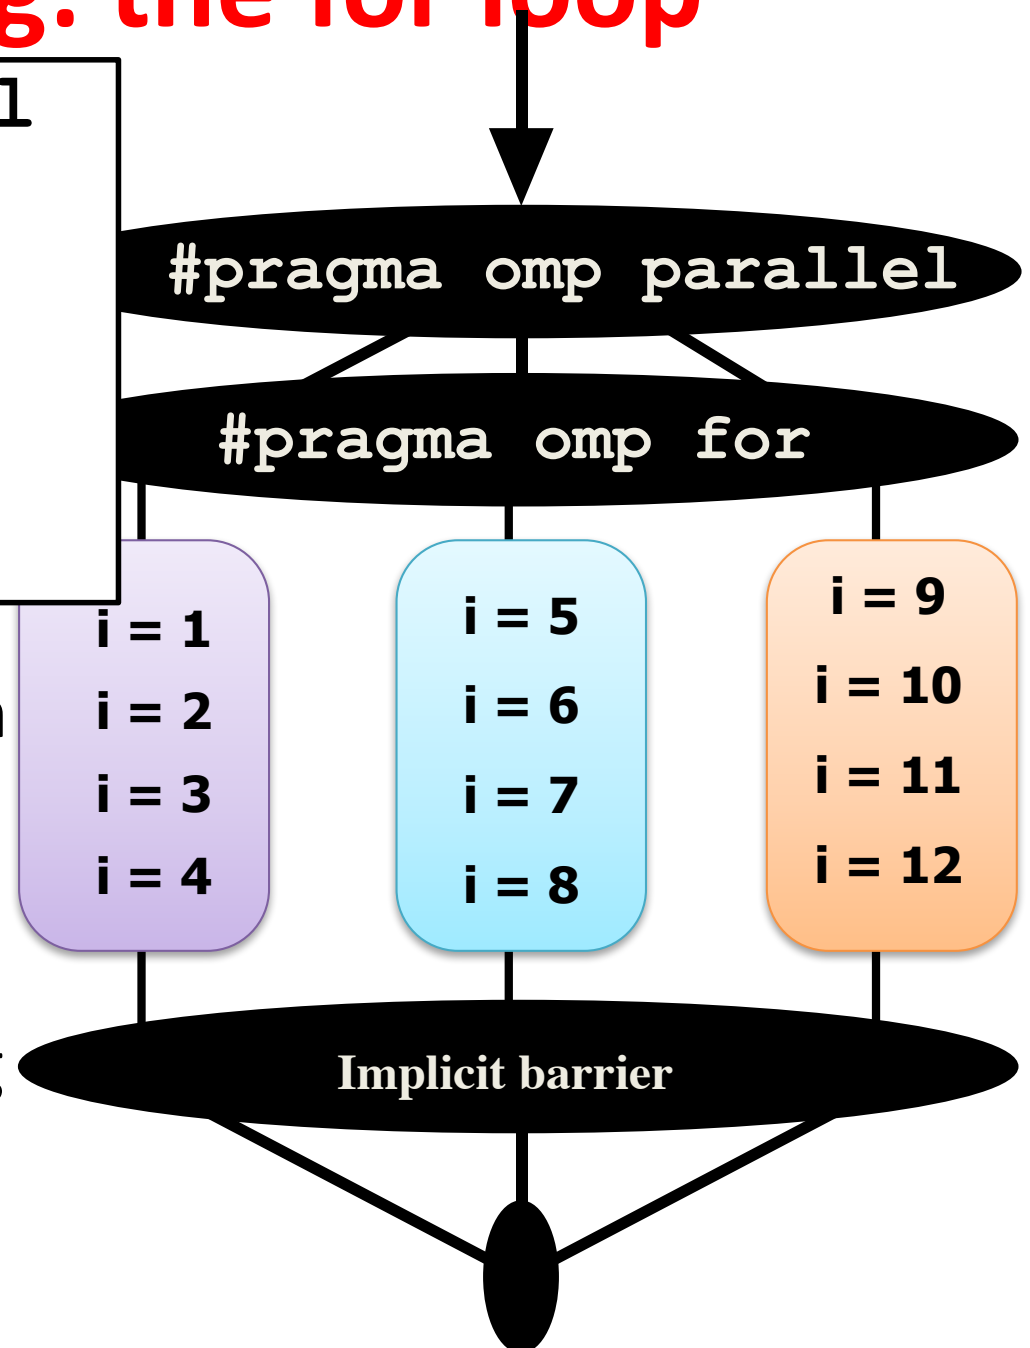
Auto parallel for loop

```
#pragma omp parallel
#pragma omp for schedule(static)
{
    for (int i=0; i<N; i++) {
        a[i]=b[i]+c[i];
    }
}
```

Work-sharing: the for loop


```
#pragma omp parallel
#pragma omp for
{
    for (i=1; i<13; i++)
        c[i]=a[i]+b[i];
}
```

- Threads are assigned an independent set of iterations
- Threads must wait at the end of work-sharing construct




OpenMP Fork-and-Join model

```
printf("begin\n");  
N = 1000;  
  
#pragma omp parallel for  
for (i=0; i<N; i++)  
    A[i] = B[i] + C[i];  
  
M = 500;  
  
#pragma omp parallel for  
for (j=0; j<M; j++)  
    p[j] = q[j] - r[j];  
printf("done\n");
```



AutoMutex: Critical Construct

```
sum = 0;
#pragma omp parallel private (lsum)
{
    lsum = 0;
    #pragma omp for
    for (i=0; i<N; i++) {
        lsum = lsum + A[i];
    }
    #pragma omp critical
    { sum += lsum; }
}
```



Threads wait their
turn;
only one thread at a
time

Reduction Clause

Shared variable



```
sum = 0;  
#pragma omp parallel for reduction (+:sum)  
{  
    for (i=0; i<N; i++) {  
        sum = sum + A[i];  
    }  
}
```

OpenMP Schedule


- Can help OpenMP decide how to handle parallelism

`schedule(type [,chunk])`

- **Schedule Types**
 - **Static** – Iterations divided into size chunk, if specified, and statically assigned to threads
 - **Dynamic** – Iterations divided into size chunk, if specified, and dynamically scheduled among threads

Static Schedule


- Although the OpenMP standard does not specify how a loop should be partitioned
- Most compilers split the loop in N/p (N #iterations, p #threads) chunks by default.
- This is called a static schedule (with chunk size N/p)
 - *For example, suppose we have a loop with 1000 iterations and 4 omp threads. The loop is partitioned as follows:*



Static Schedule with chunk

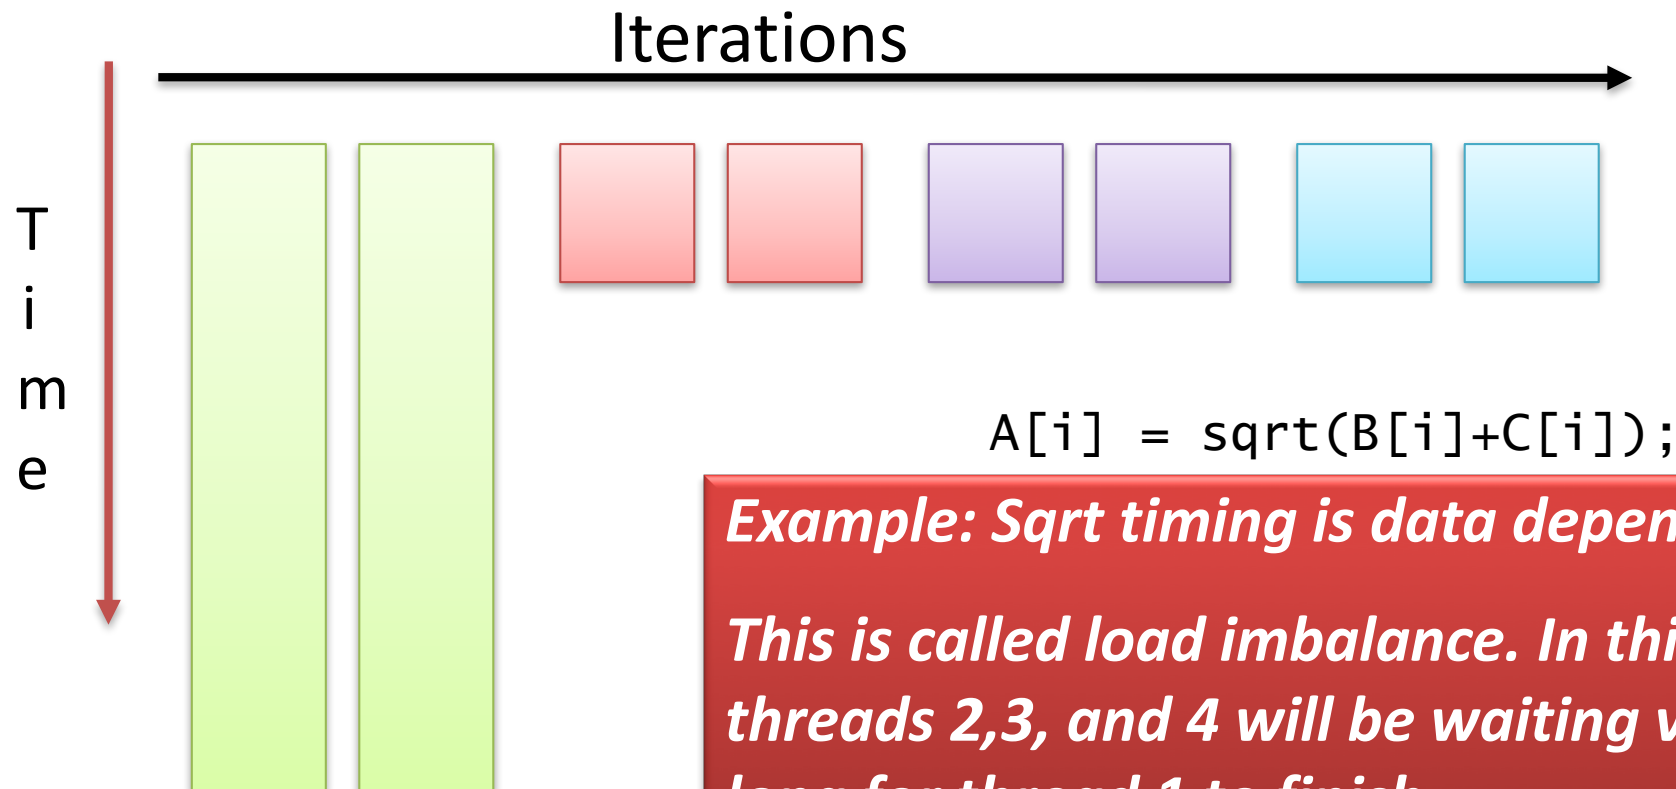
- A loop with 1000 iterations and 4 omp threads. Static Schedule with Chunk 10

```
#pragma omp parallel for schedule (static, 10)
{
  for (i=0; i<1000; i++)
    A[i] = B[i] + C[i];
}
```



Issues with Static schedule

- With static scheduling the number of iterations is evenly distributed among all openmp threads (i.e. Every thread will be assigned similar number of iterations).
- This is not always the best way to partition. Why is This?



Example: Sqrt timing is data dependent...

This is called load imbalance. In this case threads 2,3, and 4 will be waiting very long for thread 1 to finish

Dynamic Schedule

- With a dynamic schedule new chunks are assigned to threads when they come available.
- SCHEDULE(DYNAMIC,n)
 - Loop iterations are divided into pieces of size chunk. When a thread finishes one chunk, it is dynamically assigned another.

Dynamic Schedule

- SCHEDULE(GUIDED,n)
 - Similar to DYNAMIC but chunk size is relative to number of iterations left.
- Although Dynamic scheduling might be the preferred choice to prevent load imbalance
 - In some situations, there is a significant overhead involved compared to static scheduling.

More Examples on OpenMP

- <http://users.abo.fi/mats/PP2012/examples/OpenMP/>
- https://people.sc.fsu.edu/~jburkardt/c_src/mxm_openmp/mxm_openmp.html